

```
In [2]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
from sklearn.model_selection import train_test_split


import warnings
warnings.filterwarnings('ignore')
```

```
In [3]: car_data = pd.read_csv('fuel_efficiency_data.csv')
print(car_data.tail(7))
print(car_data.describe())
```

	Engine_Size_Liters	Fuel_Efficiency_MPG
93	4.757576	26.483313
94	4.797980	24.691054
95	4.838384	24.029591
96	4.878788	26.206605
97	4.919192	26.501588
98	4.959596	23.633990
99	5.000000	23.923769
	Engine_Size_Liters	Fuel_Efficiency_MPG
count	100.000000	100.000000
mean	3.000000	35.116328
std	1.172181	6.290352
min	1.000000	22.505760
25%	2.000000	29.886659
50%	3.000000	35.666563
75%	4.000000	39.902106
max	5.000000	45.990208

```
In [4]: indp_vars = car_data[['Engine_Size_Liters']]
dep_var = car_data['Fuel_Efficiency_MPG']
```

```
In [5]: plt.scatter(indp_vars, dep_var)
plt.xlabel("Engine Size (L)")
plt.ylabel("Fuel Efficiency (MPG)")
plt.title("Engine Size vs Fuel Efficiency")
plt.show()
```

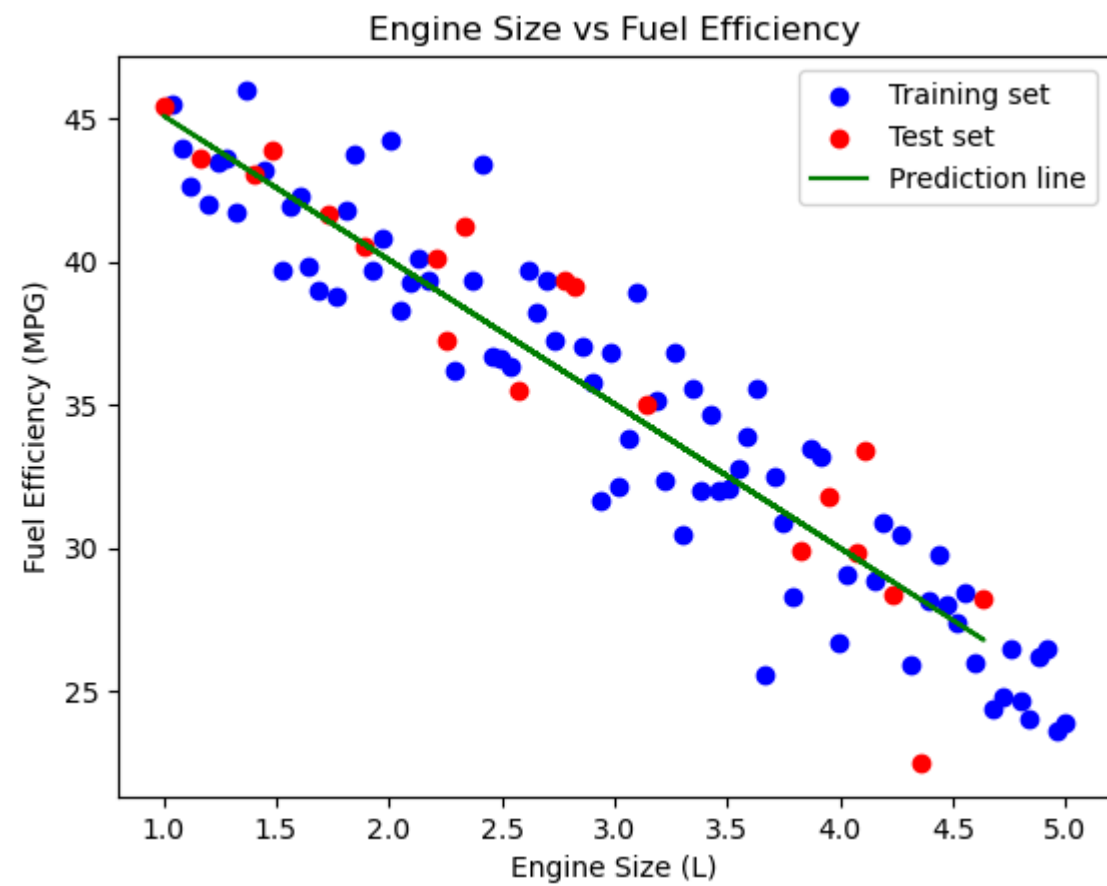


```
In [6]: train_x, test_x, train_y, test_y = train_test_split(indp_vars, dep_var, test_size=0.2, random_state=42)
lr_model = LinearRegression()
lr_model.fit(train_x, train_y)
pred = lr_model.predict(test_x)
```

```
In [7]: mse = mean_squared_error(pred, test_y)
print("Mean Squared Error:", mse)
```

Mean Squared Error: 4.570596048538351

```
In [8]: plt.scatter(train_x, train_y, color='blue', label='Training set')
plt.scatter(test_x, test_y, color='red', label='Test set')
plt.plot(test_x, pred, color="green", label="Prediction line")
plt.legend()
plt.xlabel("Engine Size (L)")
plt.ylabel("Fuel Efficiency (MPG)")
plt.title("Engine Size vs Fuel Efficiency")
plt.show()
```



```
In [10]: new_entry = input("Please enter an engine size in liters: ")
new_entry = np.array([[float(new_entry)]])
pred_new = lr_model.predict(new_entry)
print("Predicted Fuel Efficiency (MPG):", pred_new)
```

Predicted Fuel Efficiency (MPG): [35.02751255]

In [ ]:

```
In [11]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score, mean_absolute_error, mean_squared_error
```

```
In [12]: df = pd.read_csv("student_performance.csv")
```

```
In [13]: print(df.head())
print(df.info())
print(df.describe())
```

	student_id	weekly_self_study_hours	attendance_percentage	\
0	1	18.5	95.6	
1	2	14.0	80.0	
2	3	19.5	86.3	
3	4	25.7	70.2	
4	5	13.4	81.9	

	class_participation	total_score	grade
0	3.8	97.9	A
1	2.5	83.9	B
2	5.3	100.0	A
3	7.0	100.0	A
4	6.9	92.0	A

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 1000000 entries, 0 to 999999

Data columns (total 6 columns):

#	Column	Non-Null Count	Dtype
0	student_id	1000000 non-null	int64
1	weekly_self_study_hours	1000000 non-null	float64
2	attendance_percentage	1000000 non-null	float64
3	class_participation	1000000 non-null	float64
4	total_score	1000000 non-null	float64
5	grade	1000000 non-null	object

dtypes: float64(4), int64(1), object(1)


memory usage: 45.8+ MB

None

	student_id	weekly_self_study_hours	attendance_percentage	\
count	1000000.000000	1000000.000000	1000000.000000	
mean	500000.500000	15.029127	84.711046	
std	288675.278933	6.899431	9.424143	
min	1.000000	0.000000	50.000000	
25%	250000.750000	10.300000	78.300000	
50%	500000.500000	15.000000	85.000000	
75%	750000.250000	19.700000	91.800000	
max	1000000.000000	40.000000	100.000000	

	class_participation	total_score
count	1000000.000000	1000000.000000
mean	5.985203	84.283845
std	1.956421	15.432969
min	0.000000	9.400000
25%	4.700000	73.900000
50%	6.000000	87.500000
75%	7.300000	100.000000
max	10.000000	100.000000

```
In [14]: plt.figure(figsize=(6,4))
plt.scatter(df['weekly_self_study_hours'], df['total_score'])
plt.title('Weekly Self Study Hours vs Total Score')
plt.xlabel('Weekly Self Study Hours')
plt.ylabel('Total Score')
plt.grid(True)
plt.show()
```



```
In [15]: corr = df['weekly_self_study_hours'].corr(df['total_score'])
print(f"Pearson correlation: {corr:.4f}")
```


Pearson correlation: 0.8122

```
In [16]: X = df[['weekly_self_study_hours']] # independent variable
y = df['total_score'] # dependent variable
```

```
In [17]: X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.3, random_state=42
)
print("Train size:", X_train.shape, "Test size:", X_test.shape)
```

Train size: (700000, 1) Test size: (300000, 1)

```
In [18]: plt.figure(figsize=(6,4))
plt.scatter(X_train, y_train, label='Train (70%)')
plt.scatter(X_test, y_test, marker='x', label='Test (30%)')
plt.xlabel('Weekly Self Study Hours')
plt.ylabel('Total Score')
plt.title('Train vs Test Data Split')
plt.legend()
plt.grid(True)
plt.show()
```



```
In [19]: model = LinearRegression()
model.fit(X_train, y_train)

print(f"Intercept: {model.intercept_:.2f}")
print(f"Slope: {model.coef_[0]:.2f}")
print(f"Equation: total_score = {model.intercept_:.2f} + {model.coef_[0]:.2f} * weekly_self_study_hours")
```

Intercept: 56.97  
 Slope: 1.82  
 Equation: total\_score = 56.97 + 1.82 \* weekly\_self\_study\_hours

```
In [20]: y_pred = model.predict(X_test)
```

```
In [21]: r2 = r2_score(y_test, y_pred)
mae = mean_absolute_error(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)

print("Evaluation metrics on Test Set:")
print(f"R^2: {r2:.4f}")
print(f"MAE: {mae:.2f}")
print(f"MSE: {mse:.2f}")
print(f"RMSE: {rmse:.2f}")
```

Evaluation metrics on Test Set:  
 R^2: 0.6599  
 MAE: 7.16  
 MSE: 80.83  
 RMSE: 8.99

```
In [22]: results = pd.DataFrame({
    'weekly_self_study_hours': X_test['weekly_self_study_hours'],
    'Actual total_score': y_test,
    'Predicted total_score': y_pred,
    'Residual': y_test - y_pred
```

```
})
print(results.sort_values('weekly_self_study_hours'))
```

	weekly_self_study_hours	Actual total_score	Predicted total_score \
645760	0.0	51.6	56.970614
543201	0.0	53.6	56.970614
826288	0.0	48.3	56.970614
318861	0.0	60.6	56.970614
79662	0.0	38.6	56.970614
...	...	...	...
264709	40.0	100.0	129.643047
425349	40.0	100.0	129.643047
428200	40.0	100.0	129.643047
52813	40.0	100.0	129.643047
215056	40.0	100.0	129.643047

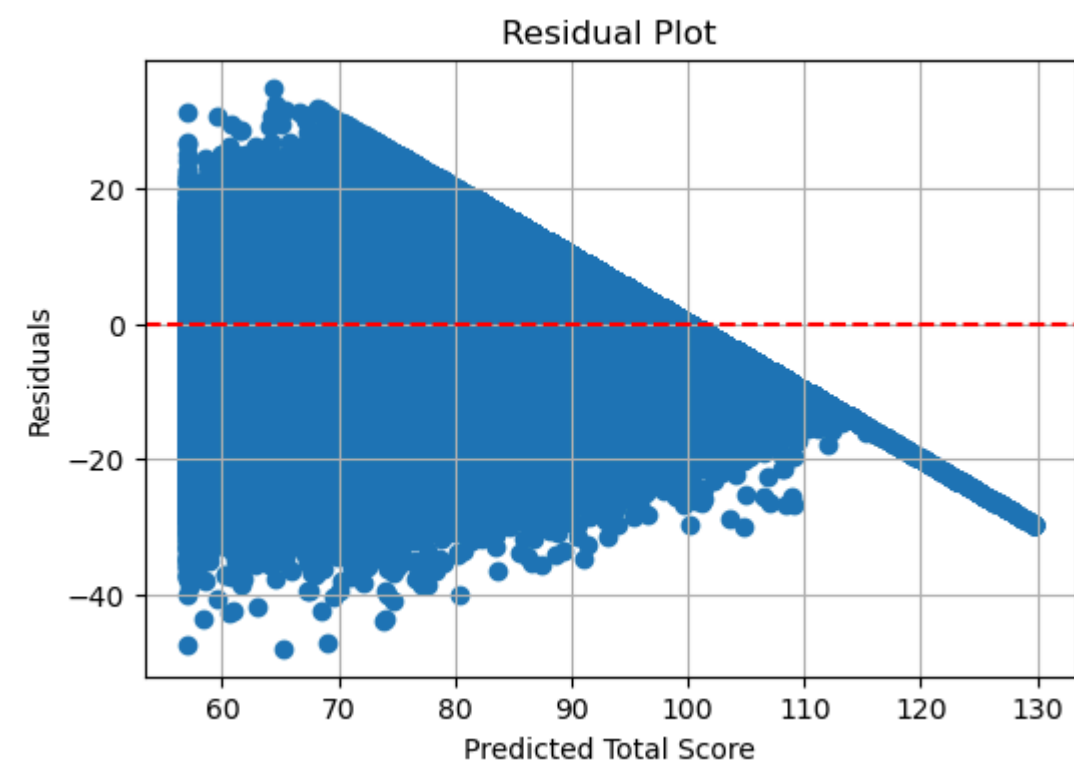
	Residual
645760	-5.370614
543201	-3.370614
826288	-8.670614
318861	3.629386
79662	-18.370614
...	...
264709	-29.643047
425349	-29.643047
428200	-29.643047
52813	-29.643047
215056	-29.643047

[300000 rows x 4 columns]

```
In [23]: plt.figure(figsize=(6,4))
# Regression Line
X_range = np.linspace(df['weekly_self_study_hours'].min(), df['weekly_self_study_hours'].max(), 200).reshape(-1,1)
plt.plot(X_range, model.predict(X_range), color='red', label='Regression Line')
# Actual vs predicted points
plt.scatter(X_test, y_test, marker='x', label='Actual (Test)')
plt.scatter(X_test, y_pred, marker='o', label='Predicted (Test)')
plt.xlabel('Weekly Self Study Hours')
plt.ylabel('Total Score')
plt.title('Regression Line with Test Data')
plt.legend()
plt.grid(True)
plt.show()
```



```
In [24]: plt.figure(figsize=(6,4))
plt.scatter(y_pred, y_test - y_pred)
plt.axhline(0, color='red', linestyle='--')
plt.xlabel('Predicted Total Score')
plt.ylabel('Residuals')
plt.title('Residual Plot')
plt.grid(True)
plt.show()
```



```
In [ ]:
```